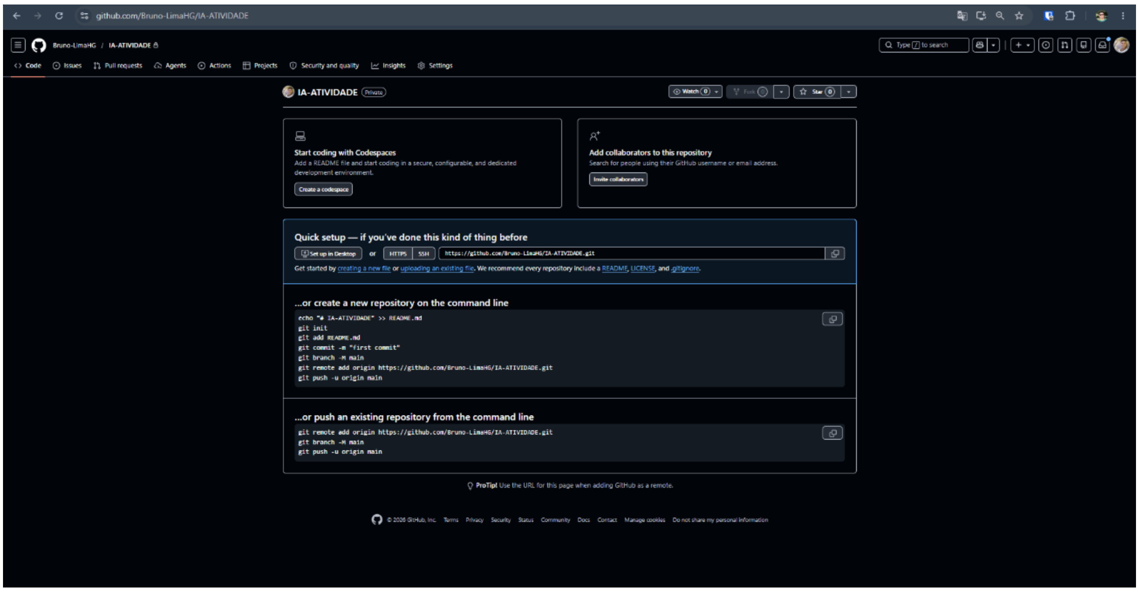
github.com/Bruno-LimaHG/IA-ATIVIDADE
Bruno-LimaHG / IA-ATIVIDADE
Code
Issues
Pull requests
Agents
Actions
Projects
Security and quality
Insights
Settings
IA-ATIVIDADE
Private
Start coding with Codespaces
Add a README file and start coding in a secure, configurable, and dedicated development environment.
Create a codespace
Add collaborators to this repository
Search for people using their GitHub username or email address.
Invite collaborators
Quick setup — if you've done this kind of thing before
Set up in Desktop
or
HTTPS
SSH
https://github.com/Bruno-LimaHG/IA-ATIVIDADE.git
Get started by creating a new file or uploading an existing file. We recommend every repository include a README, LICENSE, and .gitignore.
...or create a new repository on the command line
echo "# IA-ATIVIDADE" >> README.md
git init
git add README.md
git commit -m "first commit"
git branch -M main
git remote add origin https://github.com/Bruno-LimaHG/IA-ATIVIDADE.git
git push -u origin main
...or push an existing repository from the command line
git remote add origin https://github.com/Bruno-LimaHG/IA-ATIVIDADE.git
git branch -M main
git push -u origin main
ProTip! Use the URL for this page when adding GitHub as a remote.
© 2026 GitHub, Inc.
Terms
Privacy
Security
Status
Community
Docs
Contact
Manage cookies
Do not share my personal information

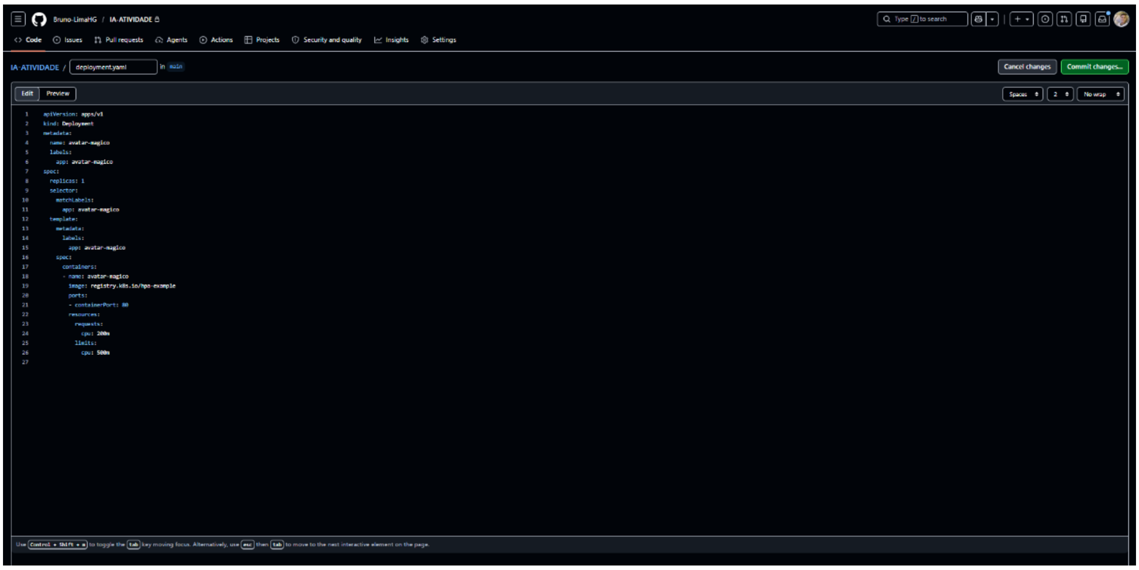
Bruno-LimaHG / IA-ATIVIDADE
Code
Issues
Pull requests
Agents
Actions
Projects
Security and quality
Insights
Settings
IA-ATIVIDADE / deployment.yaml in main
Cancel changes
Commit changes...
Edit
Preview
Spaces
2
No wrap
apiVersion: apps/v1
kind: Deployment
metadata:
name: avatar-magico
labels:
app: avatar-magico
spec:
replicas: 1
selector:
matchLabels:
app: avatar-magico
template:
metadata:
labels:
app: avatar-magico
spec:
containers:
- name: avatar-magico
image: registry.k8s.io/hpa-example
ports:
- containerPort: 80
resources:
requests:
cpu: 200m
limits:
cpu: 500m

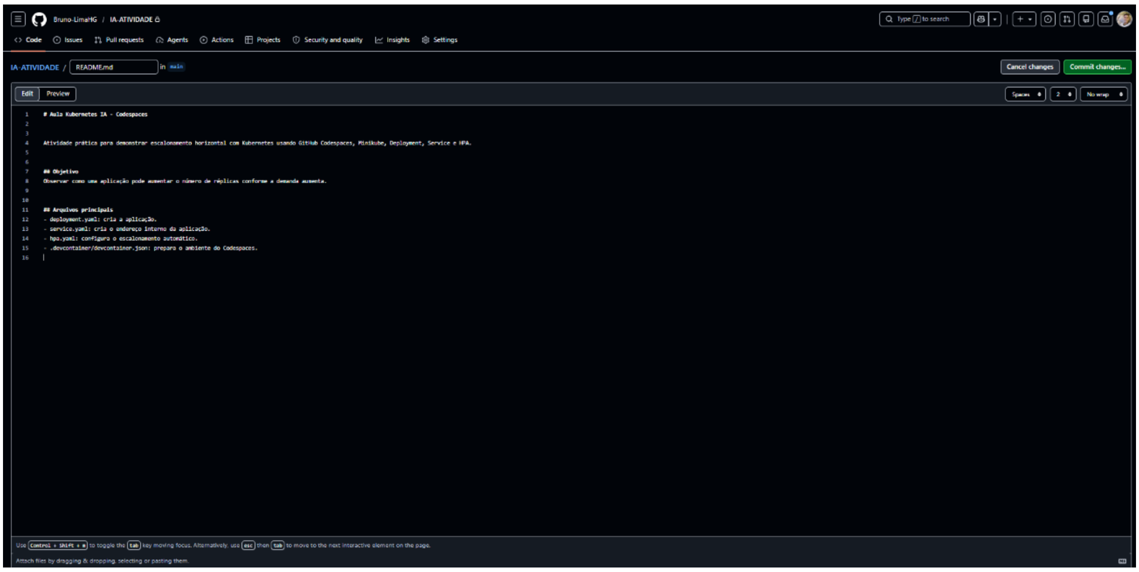
Bruno-LimaHG / IA-ATIVIDADE
Code
Issues
Pull requests
Agents
Actions
Projects
Security and quality
Insights
Settings
IA-ATIVIDADE / README.md in main
Cancel changes
Commit changes...
Edit
Preview
Spaces
2
No wrap
Aula Kubernetes IA - Codespaces
Atividade prática para demonstrar escalonamento horizontal com Kubernetes usando GitHub Codespaces, Minikube, Deployment, Service e HPA.
Objetivo
Observar como uma aplicação pode aumentar o número de réplicas conforme a demanda aumenta.
Arquivos principais
- deployment.yaml: cria a aplicação.
- service.yaml: cria o endereço interno da aplicação.
- hpa.yaml: configura o escalonamento automático.
- .devcontainer/devcontainer.json: prepara o ambiente do Codespaces.
Attach files by dragging & dropping, selecting or pasting them.

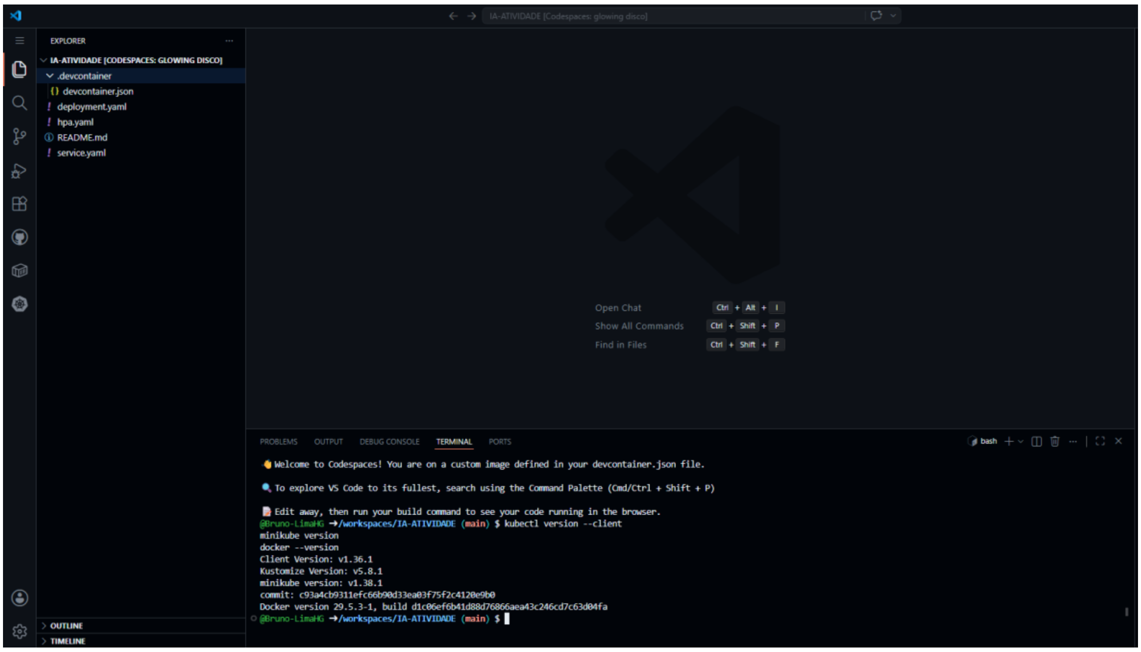
IA-ATIVIDADE [Codespaces: glowing disco]
EXPLORER
IA-ATIVIDADE [CODESPACES: GLOWING DISCO]
.devcontainer
devcontainer.json
deployment.yaml
hpa.yaml
README.md
service.yaml
Open Chat
Show All Commands
Find in Files
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
Welcome to Codespaces! You are on a custom image defined in your devcontainer.json file.
To explore VS Code to its fullest, search using the Command Palette (Cmd/Ctrl + Shift + P)
Edit away, then run your build command to see your code running in the browser.
@Bruno-LimaHG ➜ /workspaces/IA-ATIVIDADE (main) $ kubectl version --client
minikube version
docker --version
Client Version: v1.36.1
Kustomize Version: v5.8.1
minikube version: v1.38.1
commit: c93a4cb9311efc66b90d33ea03f75f2c4120e9b0
Docker version 29.5.3-1, build d1c06ef6b41d88d76866aea43c246cd7c63d04fa
@Bruno-LimaHG ➜ /workspaces/IA-ATIVIDADE (main) $
OUTLINE
TIMELINE

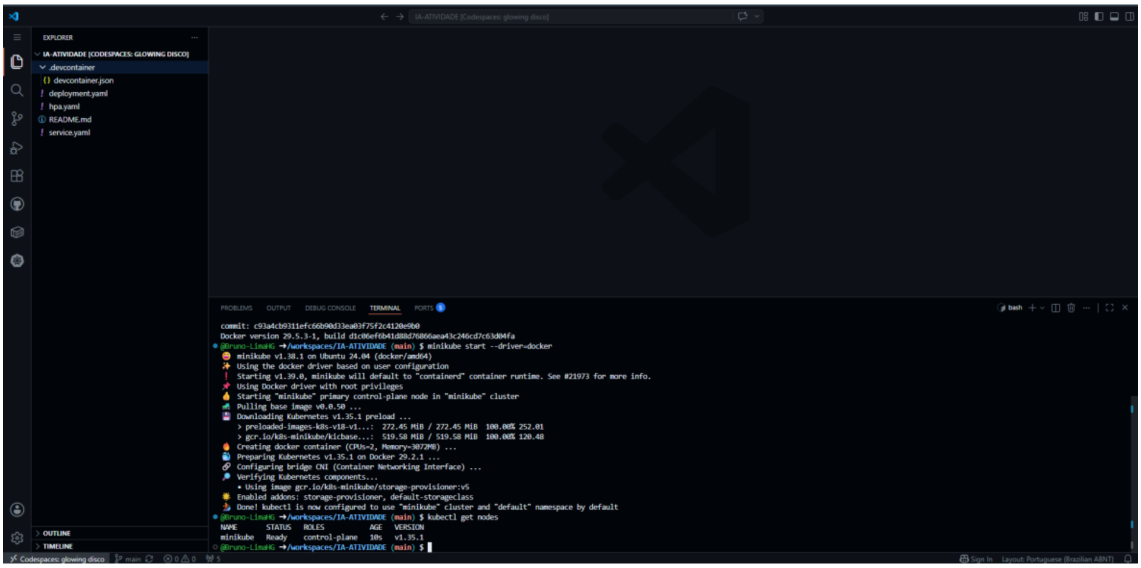
IA-ATIVIDADE [Codespaces: glowing disco]
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS 5
commit: c93a4cb9311efc66b90d33ea03f75f2c4120e9b0
Docker version 29.5.3-1, build d1c06ef6b41d88d76866aea43c246cd7c63d04fa
@Bruno-LimaHG ➜ /workspaces/IA-ATIVIDADE (main) $ minikube start --driver=docker
minikube v1.38.1 on Ubuntu 24.04 (docker/amd64)
Using the docker driver based on user configuration
Starting v1.39.0, minikube will default to "containerd" container runtime. See #21973 for more info.
Using Docker driver with root privileges
Starting "minikube" primary control-plane node in "minikube" cluster
Pulling base image v0.0.50 ...
Downloading Kubernetes v1.35.1 preload ...
> preloaded-images-k8s-v18-v1...: 272.45 MiB / 272.45 MiB 100.00% 252.01
> gcr.io/k8s-minikube/kicbase...: 519.58 MiB / 519.58 MiB 100.00% 120.48
Creating docker container (CPUs=2, Memory=3072MB) ...
Preparing Kubernetes v1.35.1 on Docker 29.2.1 ...
Configuring bridge CNI (Container Networking Interface) ...
Verifying Kubernetes components...
• Using image gcr.io/k8s-minikube/storage-provisioner:v5
Enabled addons: storage-provisioner, default-storageclass
Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
@Bruno-LimaHG ➜ /workspaces/IA-ATIVIDADE (main) $ kubectl get nodes
NAME STATUS ROLES AGE VERSION
minikube Ready control-plane 10s v1.35.1
@Bruno-LimaHG ➜ /workspaces/IA-ATIVIDADE (main) $

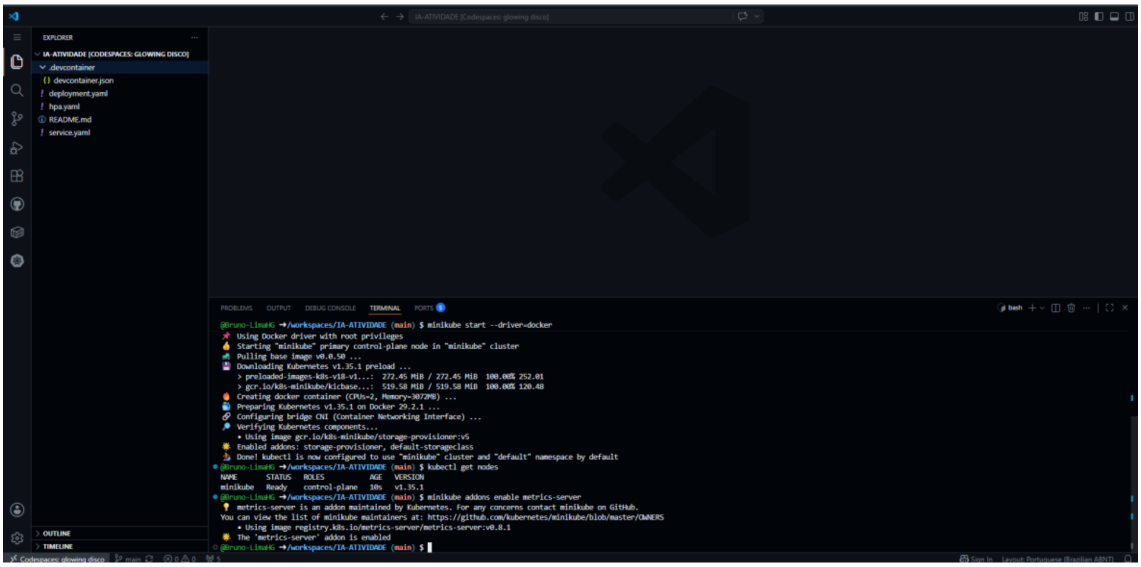
IA-ATIVIDADE [Codespaces: glowing disco]
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS 5
@Bruno-LimaHG ➜ /workspaces/IA-ATIVIDADE (main) $ minikube start --driver=docker
Using Docker driver with root privileges
Starting "minikube" primary control-plane node in "minikube" cluster
Pulling base image v0.0.50 ...
Downloading Kubernetes v1.35.1 preload ...
> preloaded-images-k8s-v18-v1...: 272.45 MiB / 272.45 MiB 100.00% 252.01
> gcr.io/k8s-minikube/kicbase...: 519.58 MiB / 519.58 MiB 100.00% 120.48
Creating docker container (CPUs=2, Memory=3072MB) ...
Preparing Kubernetes v1.35.1 on Docker 29.2.1 ...
Configuring bridge CNI (Container Networking Interface) ...
Verifying Kubernetes components...
• Using image gcr.io/k8s-minikube/storage-provisioner:v5
Enabled addons: storage-provisioner, default-storageclass
Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
@Bruno-LimaHG ➜ /workspaces/IA-ATIVIDADE (main) $ kubectl get nodes
NAME STATUS ROLES AGE VERSION
minikube Ready control-plane 10s v1.35.1
@Bruno-LimaHG ➜ /workspaces/IA-ATIVIDADE (main) $ minikube addons enable metrics-server
metrics-server is an addon maintained by Kubernetes. For any concerns contact minikube on GitHub.
You can view the list of minikube maintainers at: https://github.com/kubernetes/minikube/blob/master/OWNERS
• Using image registry.k8s.io/metrics-server/metrics-server:v0.8.1
The 'metrics-server' addon is enabled
@Bruno-LimaHG ➜ /workspaces/IA-ATIVIDADE (main) $

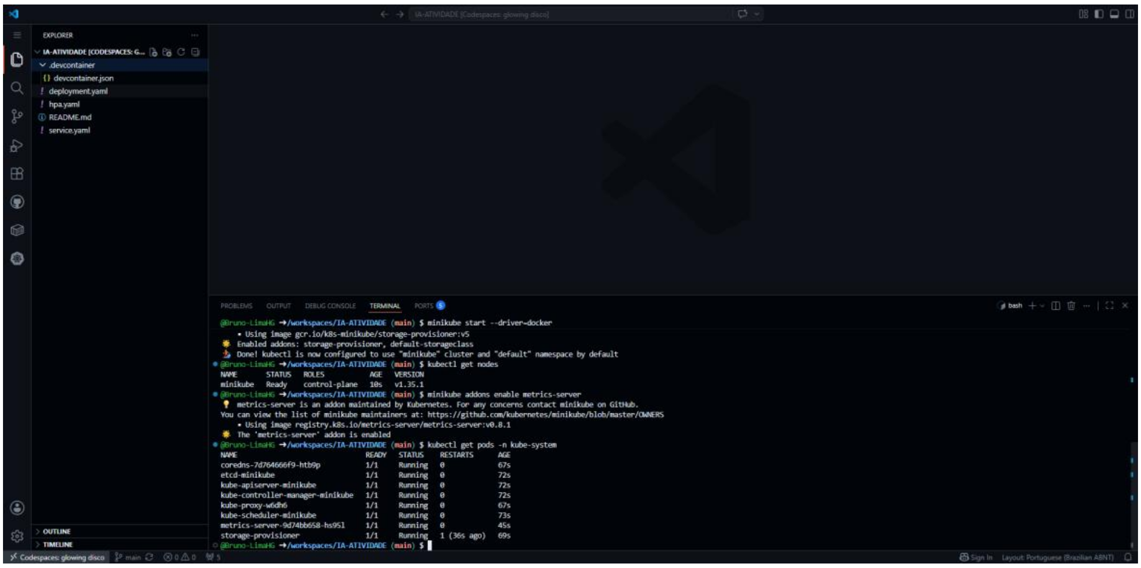
IA-ATIVIDADE [Codespaces: glowing disco]
EXPLORER
IA-ATIVIDADE [CODESPACES: G...
.devcontainer
devcontainer.json
deployment.yaml
hpa.yaml
README.md
service.yaml
OUTLINE
TIMELINE
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS 5
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ minikube start --driver=docker
• Using image gcr.io/k8s-minikube/storage-provisioner:v5
Enabled addons: storage-provisioner, default-storageclass
Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl get nodes
NAME STATUS ROLES AGE VERSION
minikube Ready control-plane 10s v1.35.1
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ minikube addons enable metrics-server
metrics-server is an addon maintained by Kubernetes. For any concerns contact minikube on GitHub.
You can view the list of minikube maintainers at: https://github.com/kubernetes/minikube/blob/master/OWNERS
• Using image registry.k8s.io/metrics-server/metrics-server:v0.8.1
The 'metrics-server' addon is enabled
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl get pods -n kube-system
NAME READY STATUS RESTARTS AGE
coredns-7d764666f9-htb9p 1/1 Running 0 67s
etcd-minikube 1/1 Running 0 72s
kube-apiserver-minikube 1/1 Running 0 72s
kube-controller-manager-minikube 1/1 Running 0 72s
kube-proxy-w6dh6 1/1 Running 0 67s
kube-scheduler-minikube 1/1 Running 0 73s
metrics-server-9d74bb658-hs95l 1/1 Running 0 45s
storage-provisioner 1/1 Running 1 (36s ago) 69s
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $
Codespaces: glowing disco
main
Sign In
Layout: Portuguese (Brazilian ABNT)

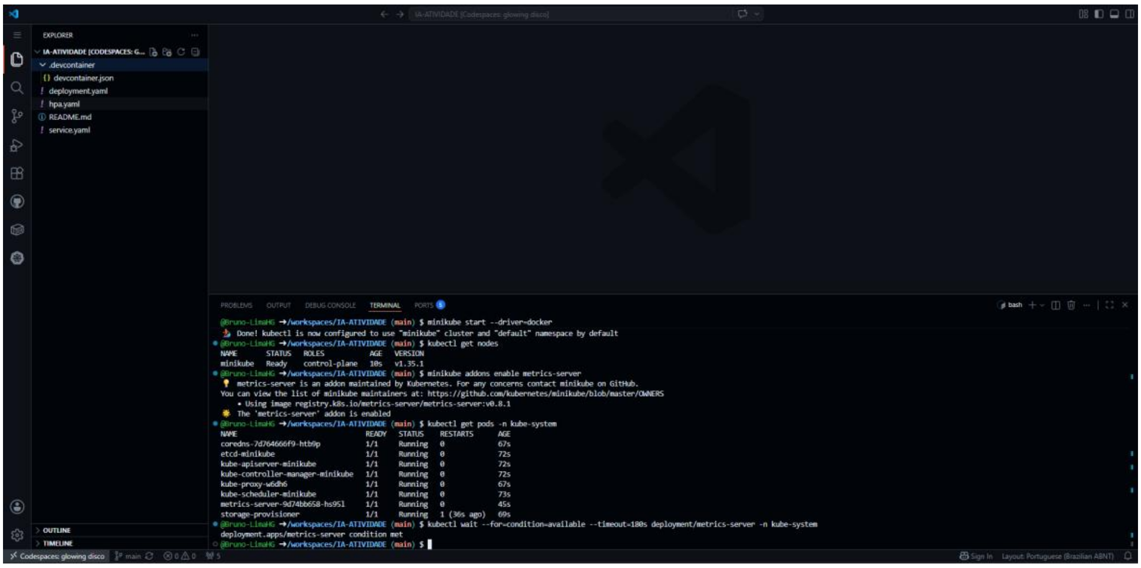
IA-ATIVIDADE [Codespaces: glowing disco]
EXPLORER
IA-ATIVIDADE [CODESPACES: G...
.devcontainer
devcontainer.json
deployment.yaml
hpa.yaml
README.md
service.yaml
OUTLINE
TIMELINE
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS 5
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ minikube start --driver=docker
Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl get nodes
NAME STATUS ROLES AGE VERSION
minikube Ready control-plane 10s v1.35.1
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ minikube addons enable metrics-server
metrics-server is an addon maintained by Kubernetes. For any concerns contact minikube on GitHub.
You can view the list of minikube maintainers at: https://github.com/kubernetes/minikube/blob/master/OWNERS
• Using image registry.k8s.io/metrics-server/metrics-server:v0.8.1
The 'metrics-server' addon is enabled
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl get pods -n kube-system
NAME READY STATUS RESTARTS AGE
coredns-7d764666f9-htb9p 1/1 Running 0 67s
etcd-minikube 1/1 Running 0 72s
kube-apiserver-minikube 1/1 Running 0 72s
kube-controller-manager-minikube 1/1 Running 0 72s
kube-proxy-w6dh6 1/1 Running 0 67s
kube-scheduler-minikube 1/1 Running 0 73s
metrics-server-9d74bb658-hs95l 1/1 Running 0 45s
storage-provisioner 1/1 Running 1 (36s ago) 69s
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl wait --for=condition=available --timeout=180s deployment/metrics-server -n kube-system
deployment.apps/metrics-server condition met
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $
Codespaces: glowing disco
main
Sign In
Layout: Portuguese (Brazilian ABNT)

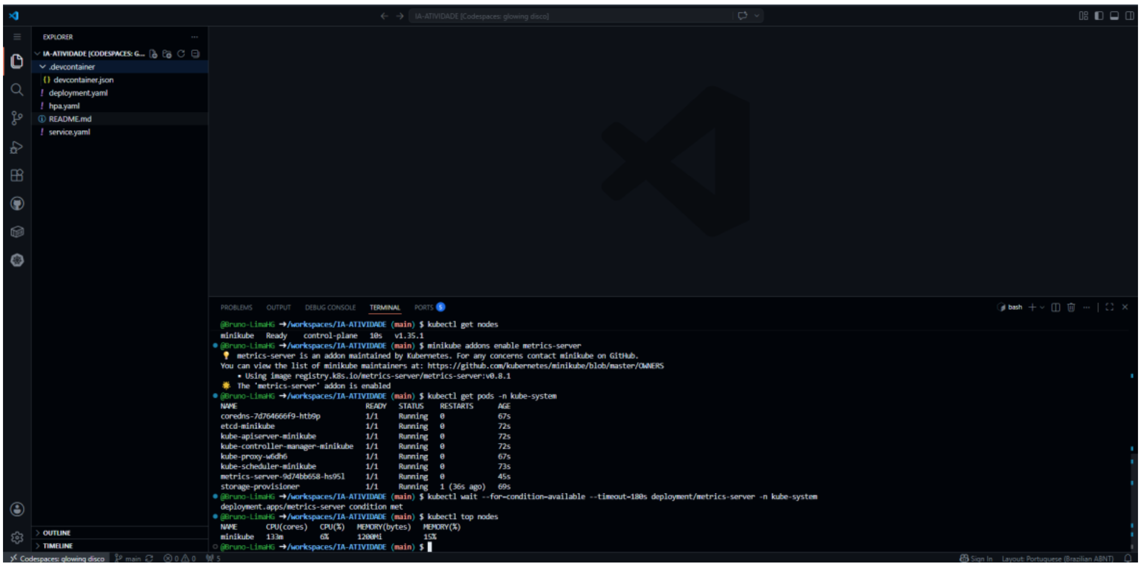
IA-ATIVIDADE [Codespaces: glowing disco]
EXPLORER
IA-ATIVIDADE [CODESPACES: G...
.devcontainer
devcontainer.json
deployment.yaml
hpa.yaml
README.md
service.yaml
OUTLINE
TIMELINE
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS 5
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl get nodes
minikube Ready control-plane 10s v1.35.1
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ minikube addons enable metrics-server
metrics-server is an addon maintained by Kubernetes. For any concerns contact minikube on GitHub.
You can view the list of minikube maintainers at: https://github.com/kubernetes/minikube/blob/master/OWNERS
• Using image registry.k8s.io/metrics-server/metrics-server:v0.8.1
The 'metrics-server' addon is enabled
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl get pods -n kube-system
NAME READY STATUS RESTARTS AGE
coredns-7d764666f9-htb9p 1/1 Running 0 67s
etcd-minikube 1/1 Running 0 72s
kube-apiserver-minikube 1/1 Running 0 72s
kube-controller-manager-minikube 1/1 Running 0 72s
kube-proxy-w6dh6 1/1 Running 0 67s
kube-scheduler-minikube 1/1 Running 0 73s
metrics-server-9d74bb658-hs95l 1/1 Running 0 45s
storage-provisioner 1/1 Running 1 (36s ago) 69s
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl wait --for=condition=available --timeout=180s deployment/metrics-server -n kube-system
deployment.apps/metrics-server condition met
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl top nodes
NAME CPU(cores) CPU(%) MEMORY(bytes) MEMORY(%)
minikube 133m 6% 1200Mi 15%
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $
Codespaces: glowing disco
main
Sign In
Layout: Portuguese (Brazilian ABNT)

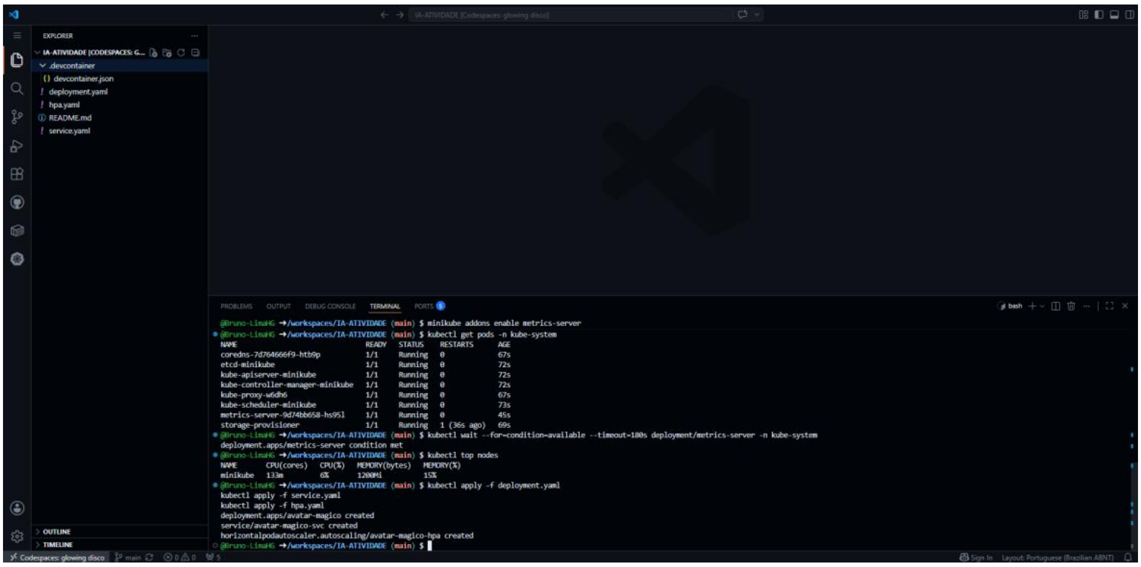
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $ minikube addons enable metrics-server
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $ kubectl get pods -n kube-system
NAME READY STATUS RESTARTS AGE
coredns-7d764666f9-htb9p 1/1 Running 0 67s
etcd-minikube 1/1 Running 0 72s
kube-apiserver-minikube 1/1 Running 0 72s
kube-controller-manager-minikube 1/1 Running 0 72s
kube-proxy-w6dh6 1/1 Running 0 67s
kube-scheduler-minikube 1/1 Running 0 73s
metrics-server-9d74bb658-hs951 1/1 Running 0 45s
storage-provisioner 1/1 Running 1 (36s ago) 69s
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $ kubectl wait --for=condition=available --timeout=180s deployment/metrics-server -n kube-system
deployment.apps/metrics-server condition met
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $ kubectl top nodes
NAME CPU(cores) CPU(%) MEMORY(bytes) MEMORY(%)
minikube 133m 6% 1200Mi 15%
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $ kubectl apply -f deployment.yaml
kubectl apply -f service.yaml
kubectl apply -f hpa.yaml
deployment.apps/avatar-magico created
service/avatar-magico-svc created
horizontalpodautoscaler.autoscaling/avatar-magico-hpa created
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $

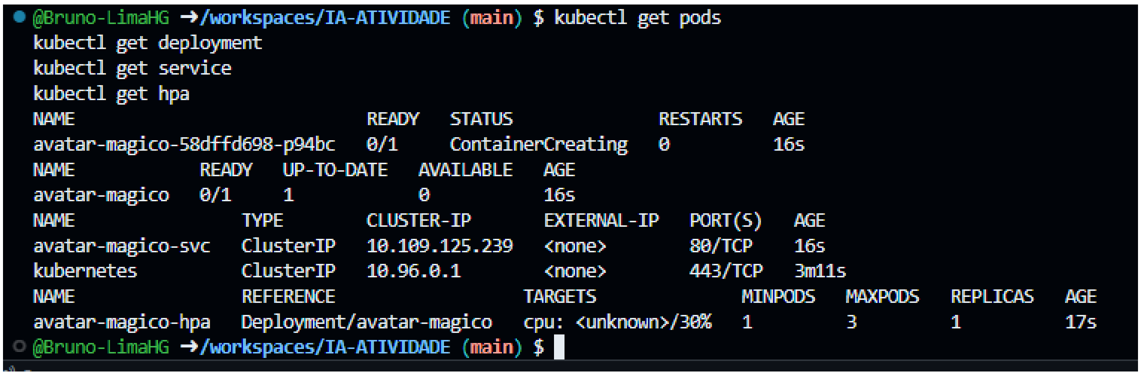
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $ kubectl get pods
kubectl get deployment
kubectl get service
kubectl get hpa
NAME READY STATUS RESTARTS AGE
avatar-magico-58dffd698-p94bc 0/1 ContainerCreating 0 16s
NAME READY UP-TO-DATE AVAILABLE AGE
avatar-magico 0/1 1 0 16s
NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
avatar-magico-svc ClusterIP 10.109.125.239 <none> 80/TCP 16s
kubernetes ClusterIP 10.96.0.1 <none> 443/TCP 3m11s
NAME REFERENCE TARGETS MINPODS MAXPODS REPLICAS AGE
avatar-magico-hpa Deployment/avatar-magico cpu: <unknown>/30% 1 3 1 17s
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $

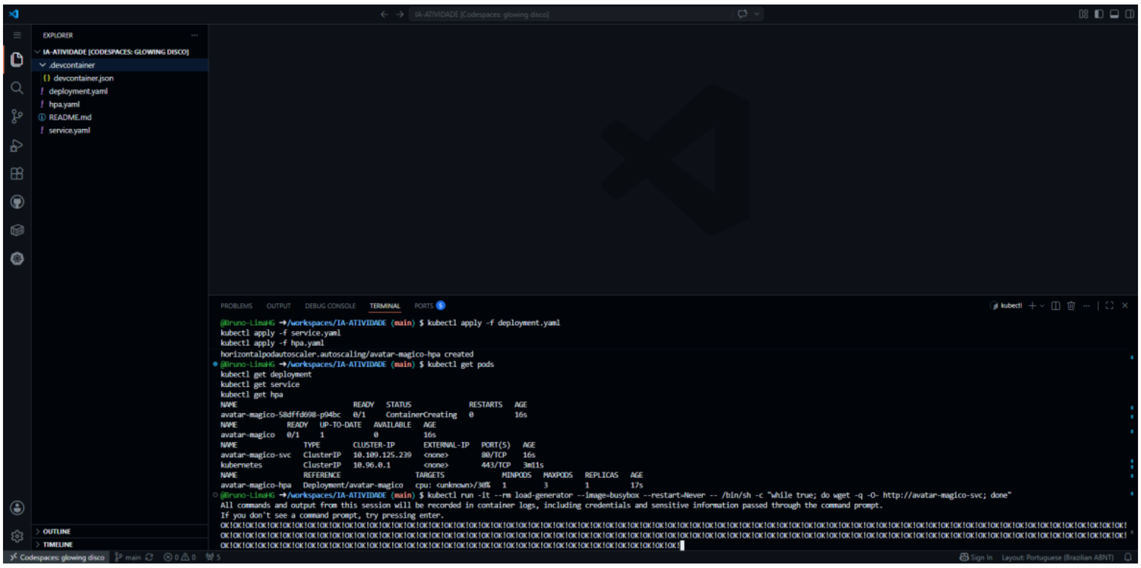
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $ kubectl apply -f deployment.yaml
kubectl apply -f service.yaml
kubectl apply -f hpa.yaml
horizontalpodautoscaler.autoscaling/avatar-magico-hpa created
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $ kubectl get pods
kubectl get deployment
kubectl get service
kubectl get hpa
NAME READY STATUS RESTARTS AGE
avatar-magico-58dffd698-p94bc 0/1 ContainerCreating 0 16s
NAME READY UP-TO-DATE AVAILABLE AGE
avatar-magico 0/1 1 0 16s
NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
avatar-magico-svc ClusterIP 10.109.125.239 <none> 80/TCP 16s
kubernetes ClusterIP 10.96.0.1 <none> 443/TCP 3m11s
NAME REFERENCE TARGETS MINPODS MAXPODS REPLICAS AGE
avatar-magico-hpa Deployment/avatar-magico cpu: <unknown>/30% 1 3 1 17s
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $ kubectl run -it --rm load-generator --image=busybox --restart=Never -- /bin/sh -c "while true; do wget -q -O- http://avatar-magico-svc; done"
All commands and output from this session will be recorded in container logs, including credentials and sensitive information passed through the command prompt.
If you don't see a command prompt, try pressing enter.

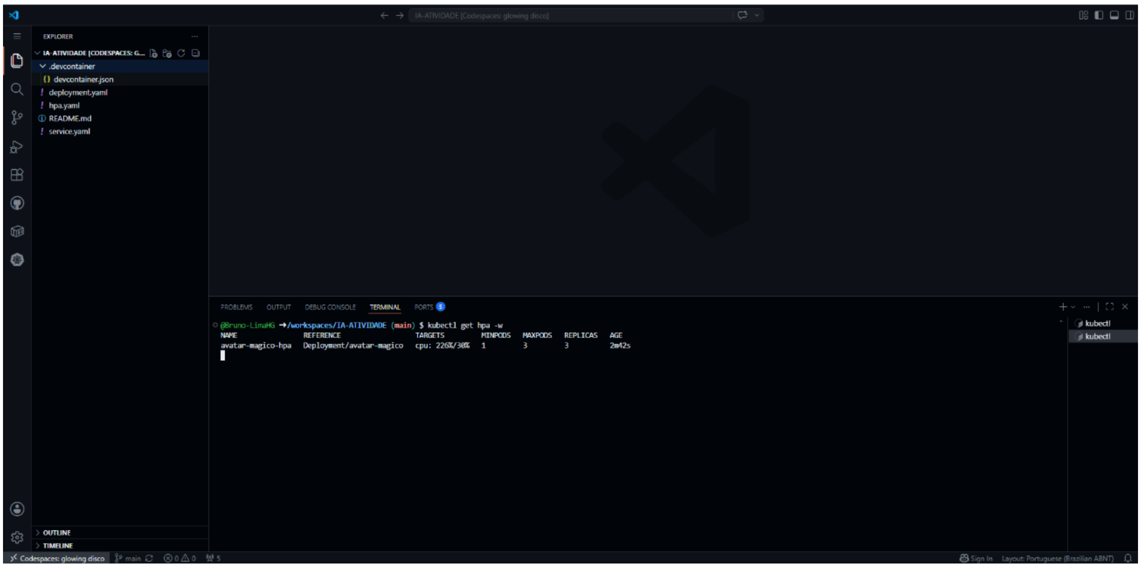
EXPLORER
IA-ATIVIDADE [CODESPACES: G...
.devcontainer
devcontainer.json
deployment.yaml
hpa.yaml
README.md
service.yaml
PROBLEMS
OUTPUT
DEBUG CONSOLE
TERMINAL
PORTS
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl get hpa -w
NAME REFERENCE TARGETS MINPODS MAXPODS REPLICAS AGE
avatar-magico-hpa Deployment/avatar-magico cpu: 226%/30% 1 3 3 2m42s
OUTLINE
TIMELINE

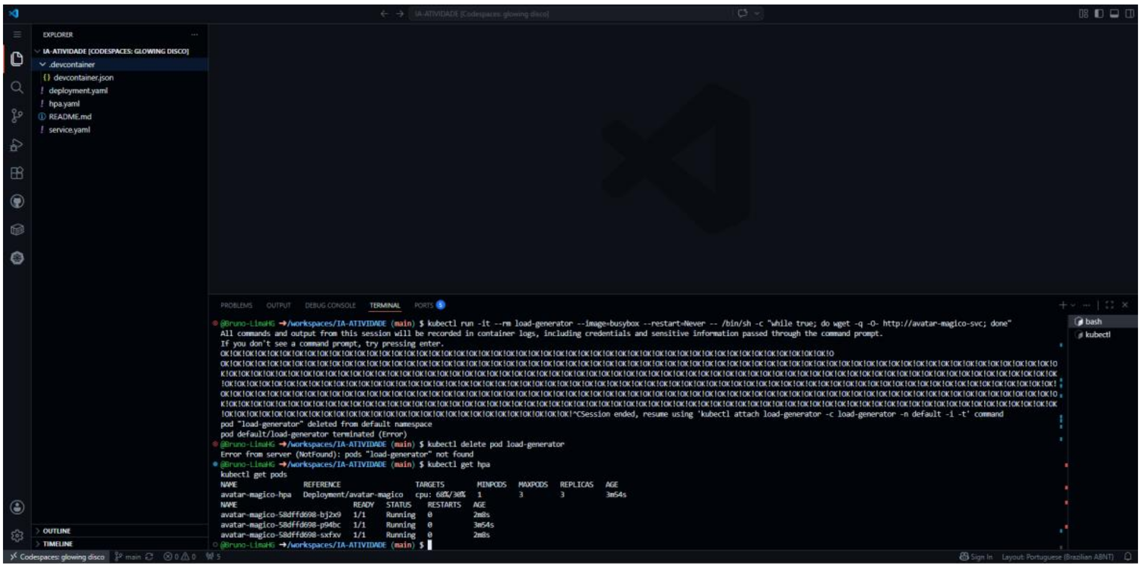
EXPLORER
IA-ATIVIDADE [CODESPACES: GLOWING DISCO]
.devcontainer
devcontainer.json
deployment.yaml
hpa.yaml
README.md
service.yaml
PROBLEMS
OUTPUT
DEBUG CONSOLE
TERMINAL
PORTS
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl run -it --rm load-generator --image=busybox --restart=Never -- /bin/sh -c "while true; do wget -q -O- http://avatar-magico-svc; done"
All commands and output from this session will be recorded in container logs, including credentials and sensitive information passed through the command prompt.
If you don't see a command prompt, try pressing enter.
^CSession ended, resume using 'kubectl attach load-generator -c load-generator -n default -i -t' command
pod "load-generator" deleted from default namespace
pod default/load-generator terminated (Error)
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl delete pod load-generator
Error from server (NotFound): pods "load-generator" not found
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl get hpa
kubectl get pods
NAME REFERENCE TARGETS MINPODS MAXPODS REPLICAS AGE
avatar-magico-hpa Deployment/avatar-magico cpu: 68%/30% 1 3 3 3m54s
NAME READY STATUS RESTARTS AGE
avatar-magico-58dffd698-bj2x9 1/1 Running 0 2m8s
avatar-magico-58dffd698-p94bc 1/1 Running 0 3m54s
avatar-magico-58dffd698-sxfxv 1/1 Running 0 2m8s
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $
OUTLINE
TIMELINE

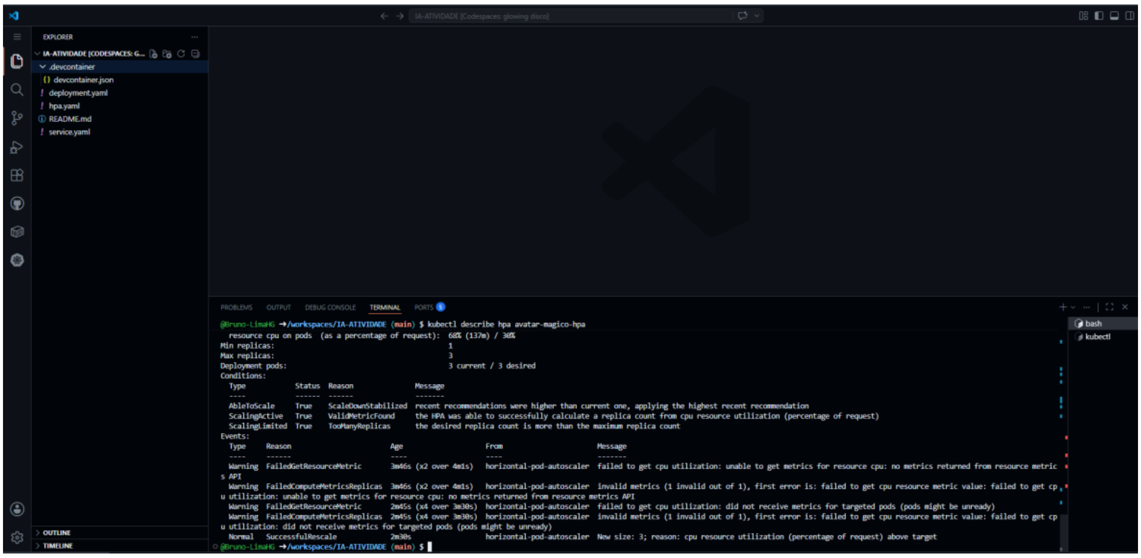
EXPLORER
IA-ATIVIDADE [CODESPACES: G...
.devcontainer
devcontainer.json
deployment.yaml
hpa.yaml
README.md
service.yaml
PROBLEMS
OUTPUT
DEBUG CONSOLE
TERMINAL
PORTS
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl describe hpa avatar-magico-hpa
resource cpu on pods (as a percentage of request): 68% (137m) / 30%
Min replicas: 1
Max replicas: 3
Deployment pods: 3 current / 3 desired
Conditions:
Type Status Reason Message
AbleToScale True ScaleDownStabilized recent recommendations were higher than current one, applying the highest recent recommendation
ScalingActive True ValidMetricFound the HPA was able to successfully calculate a replica count from cpu resource utilization (percentage of request)
ScalingLimited True TooManyReplicas the desired replica count is more than the maximum replica count
Events:
Type Reason Age From Message
Warning FailedGetResourceMetric 3m46s (x2 over 4m1s) horizontal-pod-autoscaler failed to get cpu utilization: unable to get metrics for resource cpu: no metrics returned from resource metrics API
Warning FailedComputeMetricsReplicas 3m46s (x2 over 4m1s) horizontal-pod-autoscaler invalid metrics (1 invalid out of 1), first error is: failed to get cpu resource metric value: failed to get cpu utilization: unable to get metrics for resource cpu: no metrics returned from resource metrics API
Warning FailedGetResourceMetric 2m45s (x4 over 3m30s) horizontal-pod-autoscaler failed to get cpu utilization: did not receive metrics for targeted pods (pods might be unready)
Warning FailedComputeMetricsReplicas 2m45s (x4 over 3m30s) horizontal-pod-autoscaler invalid metrics (1 invalid out of 1), first error is: failed to get cpu resource metric value: failed to get cpu utilization: did not receive metrics for targeted pods (pods might be unready)
Normal SuccessfulRescale 2m30s horizontal-pod-autoscaler New size: 3; reason: cpu resource utilization (percentage of request) above target
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $
OUTLINE
TIMELINE

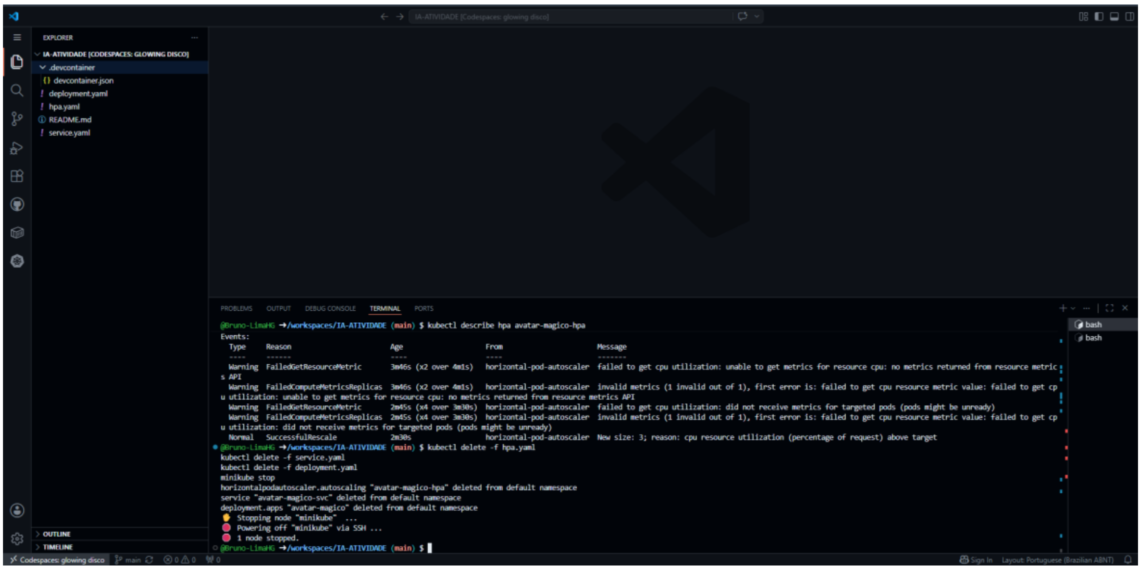
EXPLORER
IA-ATIVIDADE [CODESPACES: GLOWING DISCO]
.devcontainer
devcontainer.json
deployment.yaml
hpa.yaml
README.md
service.yaml
PROBLEMS
OUTPUT
DEBUG CONSOLE
TERMINAL
PORTS
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $ kubectl describe hpa avatar-magico-hpa
Events:
Type Reason Age From Message
Warning FailedGetResourceMetric 3m46s (x2 over 4m1s) horizontal-pod-autoscaler failed to get cpu utilization: unable to get metrics for resource cpu: no metrics returned from resource metrics API
Warning FailedComputeMetricsReplicas 3m46s (x2 over 4m1s) horizontal-pod-autoscaler invalid metrics (1 invalid out of 1), first error is: failed to get cpu resource metric value: failed to get cpu utilization: unable to get metrics for resource cpu: no metrics returned from resource metrics API
Warning FailedGetResourceMetric 2m45s (x4 over 3m30s) horizontal-pod-autoscaler failed to get cpu utilization: did not receive metrics for targeted pods (pods might be unready)
Warning FailedComputeMetricsReplicas 2m45s (x4 over 3m30s) horizontal-pod-autoscaler invalid metrics (1 invalid out of 1), first error is: failed to get cpu resource metric value: failed to get cpu utilization: did not receive metrics for targeted pods (pods might be unready)
Normal SuccessfulRescale 2m30s horizontal-pod-autoscaler New size: 3; reason: cpu resource utilization (percentage of request) above target
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $ kubectl delete -f hpa.yaml
kubectl delete -f service.yaml
kubectl delete -f deployment.yaml
minikube stop
horizontalpodautoscaler.autoscaling "avatar-magico-hpa" deleted from default namespace
service "avatar-magico-svc" deleted from default namespace
deployment.apps "avatar-magico" deleted from default namespace
Stopping node "minikube" ...
Powering off "minikube" via SSH ...
1 node stopped.
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $
Sign In
Layout: Portuguese (Brazilian ABNT)

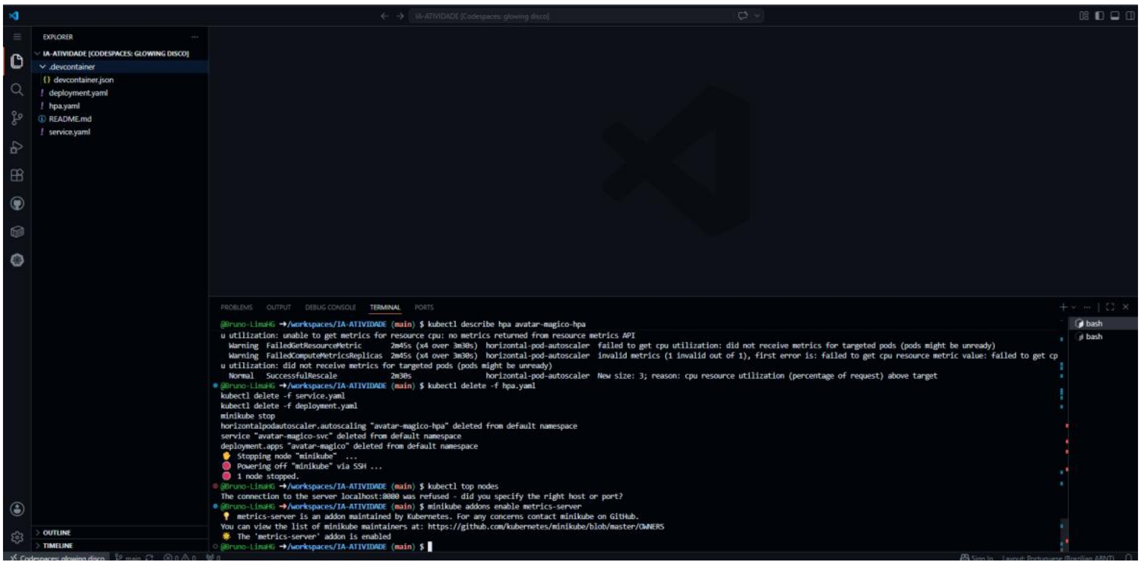
EXPLORER
IA-ATIVIDADE [CODESPACES: GLOWING DISCO]
.devcontainer
devcontainer.json
deployment.yaml
hpa.yaml
README.md
service.yaml
TERMINAL
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $ kubectl describe hpa avatar-magico-hpa
u utilization: unable to get metrics for resource cpu: no metrics returned from resource metrics API
Warning FailedGetResourceMetric 2m45s (x4 over 3m30s) horizontal-pod-autoscaler failed to get cpu utilization: did not receive metrics for targeted pods (pods might be unready)
Warning FailedComputeMetricsReplicas 2m45s (x4 over 3m30s) horizontal-pod-autoscaler invalid metrics (1 invalid out of 1), first error is: failed to get cpu resource metric value: failed to get cpu utilization: did not receive metrics for targeted pods (pods might be unready)
Normal SuccessfulRescale 2m30s horizontal-pod-autoscaler New size: 3; reason: cpu resource utilization (percentage of request) above target
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $ kubectl delete -f hpa.yaml
kubectl delete -f service.yaml
kubectl delete -f deployment.yaml
minikube stop
horizontalpodautoscaler.autoscaling "avatar-magico-hpa" deleted from default namespace
service "avatar-magico-svc" deleted from default namespace
deployment.apps "avatar-magico" deleted from default namespace
Stopping node "minikube" ...
Powering off "minikube" via SSH ...
1 node stopped.
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $ kubectl top nodes
The connection to the server localhost:8080 was refused - did you specify the right host or port?
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $ minikube addons enable metrics-server
metrics-server is an addon maintained by Kubernetes. For any concerns contact minikube on GitHub.
You can view the list of minikube maintainers at: https://github.com/kubernetes/minikube/blob/master/OWNERS
The 'metrics-server' addon is enabled
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $

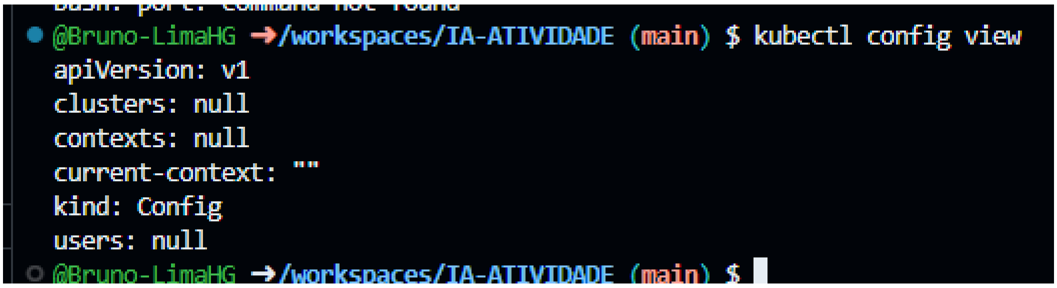
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $ kubectl config view
apiVersion: v1
clusters: null
contexts: null
current-context: ""
kind: Config
users: null
@Bruno-LimaHG →/workspaces/IA-ATIVIDADE (main) $

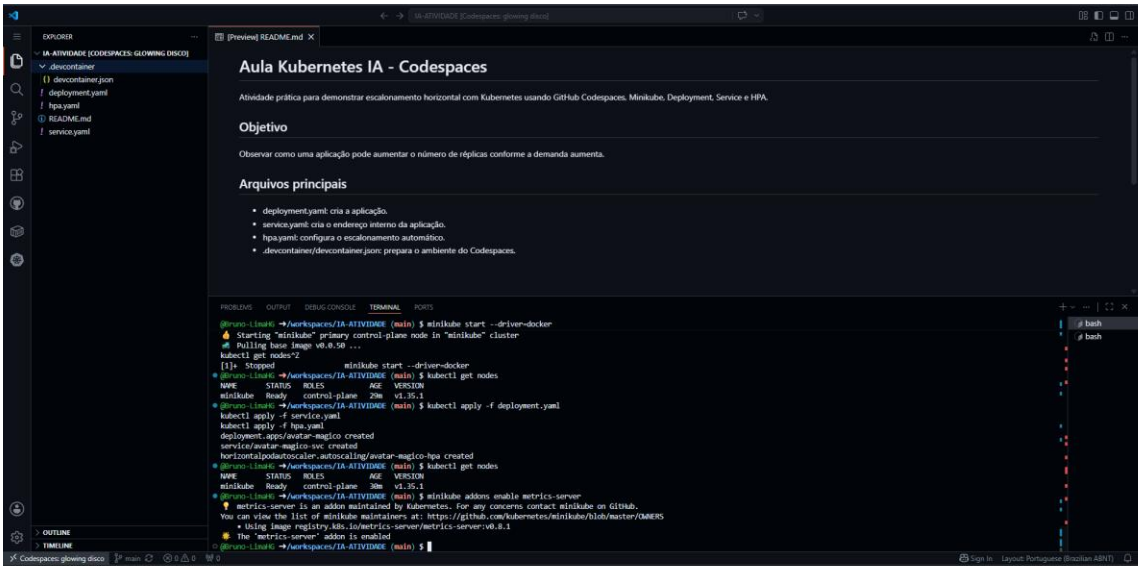
IA-ATIVIDADE [Codespaces: glowing disco]
EXPLORER
IA-ATIVIDADE [CODESPACES: GLOWING DISCO]
.devcontainer
devcontainer.json
deployment.yaml
hpa.yaml
README.md
service.yaml
OUTLINE
TIMELINE
[Preview] README.md
Aula Kubernetes IA - Codespaces
Atividade prática para demonstrar escalonamento horizontal com Kubernetes usando GitHub Codespaces, Minikube, Deployment, Service e HPA.
Objetivo
Observar como uma aplicação pode aumentar o número de réplicas conforme a demanda aumenta.
Arquivos principais
deployment.yaml: cria a aplicação.
service.yaml: cria o endereço interno da aplicação.
hpa.yaml: configura o escalonamento automático.
.devcontainer/devcontainer.json: prepara o ambiente do Codespaces.
PROBLEMS
OUTPUT
DEBUG CONSOLE
TERMINAL
PORTS
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ minikube start --driver=docker
Starting "minikube" primary control-plane node in "minikube" cluster
Pulling base image v0.0.50 ...
kubectl get nodes^Z
[1]+ Stopped   minikube start --driver=docker
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl get nodes
NAME   STATUS   ROLES   AGE   VERSION
minikube   Ready   control-plane   29m   v1.35.1
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl apply -f deployment.yaml
kubectl apply -f service.yaml
kubectl apply -f hpa.yaml
deployment.apps/avatar-magico created
service/avatar-magico-svc created
horizontalpodautoscaler.autoscaling/avatar-magico-hpa created
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl get nodes
NAME   STATUS   ROLES   AGE   VERSION
minikube   Ready   control-plane   30m   v1.35.1
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ minikube addons enable metrics-server
metrics-server is an addon maintained by Kubernetes. For any concerns contact minikube on GitHub.
You can view the list of minikube maintainers at: https://github.com/kubernetes/minikube/blob/master/OWNERS
• Using image registry.k8s.io/metrics-server/metrics-server:v0.8.1
The 'metrics-server' addon is enabled
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $
bash
bash
Codespaces: glowing disco
main
Sign In
Layout: Portuguese (Brazilian ABNT)

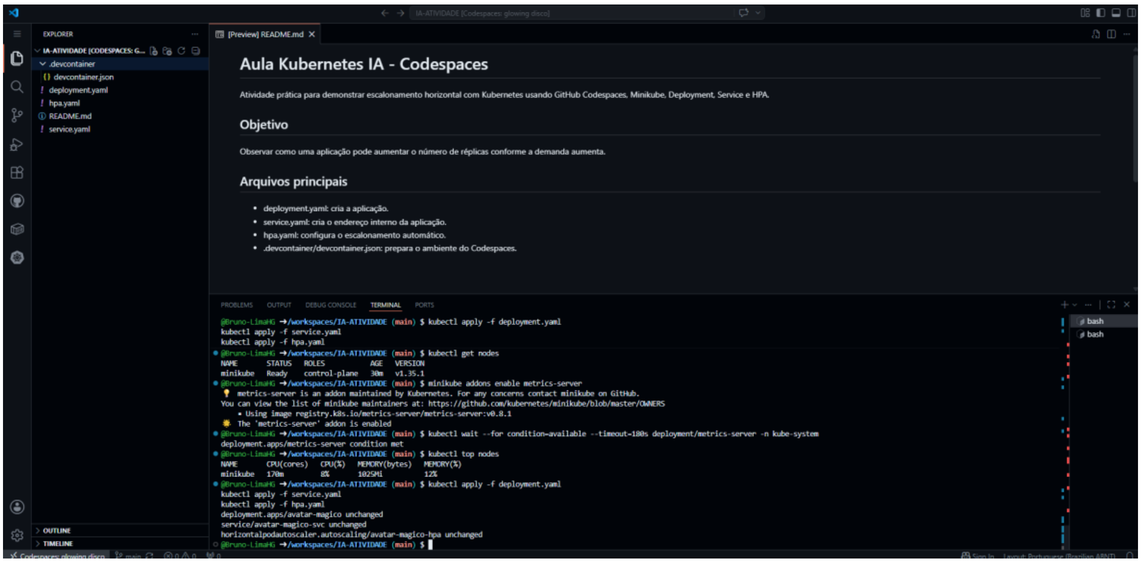
IA-ATIVIDADE [Codespaces: glowing disco]
EXPLORER
.devcontainer
devcontainer.json
deployment.yaml
hpa.yaml
README.md
service.yaml
OUTLINE
TIMELINE
[Preview] README.md
Aula Kubernetes IA - Codespaces
Atividade prática para demonstrar escalonamento horizontal com Kubernetes usando GitHub Codespaces, Minikube, Deployment, Service e HPA.
Objetivo
Observar como uma aplicação pode aumentar o número de réplicas conforme a demanda aumenta.
Arquivos principais
deployment.yaml: cria a aplicação.
service.yaml: cria o endereço interno da aplicação.
hpa.yaml: configura o escalonamento automático.
.devcontainer/devcontainer.json: prepara o ambiente do Codespaces.
PROBLEMS
OUTPUT
DEBUG CONSOLE
TERMINAL
PORTS
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl apply -f deployment.yaml
kubectl apply -f service.yaml
kubectl apply -f hpa.yaml
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl get nodes
NAME   STATUS   ROLES   AGE   VERSION
minikube   Ready   control-plane   30m   v1.35.1
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ minikube addons enable metrics-server
metrics-server is an addon maintained by Kubernetes. For any concerns contact minikube on GitHub.
You can view the list of minikube maintainers at: https://github.com/kubernetes/minikube/blob/master/OWNERS
• Using image registry.k8s.io/metrics-server/metrics-server:v0.8.1
The 'metrics-server' addon is enabled
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl wait --for condition=available --timeout=180s deployment/metrics-server -n kube-system
deployment.apps/metrics-server condition met
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl top nodes
NAME   CPU(cores)   CPU(%)   MEMORY(bytes)   MEMORY(%)
minikube   170m   8%   1025Mi   12%
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $ kubectl apply -f deployment.yaml
kubectl apply -f service.yaml
kubectl apply -f hpa.yaml
deployment.apps/avatar-magico unchanged
service/avatar-magico-svc unchanged
horizontalpodautoscaler.autoscaling/avatar-magico-hpa unchanged
@Bruno-LimaHG ➜/workspaces/IA-ATIVIDADE (main) $
bash
bash